



Tribhuvan University
Institute of Engineering
Purwanchal Campus
पूर्वाञ्चल क्याम्पस

Unit 6

Data Engineering and Automation

Contents

- 6.1 Overview of data engineering in applied data science
- 6.2 Designing and implementing ETL pipelines
- 6.3 Automating workflows with schedulers (CRON, schedule)
- 6.4 Logging, monitoring, and error handling in pipelines
- 6.5 Data storage and retrieval strategies for pipelines
- 6.6 Automated report generation (Excel, HTML, PDF)
- 6.7 Case study: End-to-end automated analytics pipeline

Contents

- 6.1 Overview of data engineering in applied data science**
- 6.2 Designing and implementing ETL pipelines
- 6.3 Automating workflows with schedulers (CRON, schedule)
- 6.4 Logging, monitoring, and error handling in pipelines
- 6.5 Data storage and retrieval strategies for pipelines
- 6.6 Automated report generation (Excel, HTML, PDF)
- 6.7 Case study: End-to-end automated analytics pipeline

Data Engineering

- Data engineering is the discipline that designs, builds, and maintains systems for collecting, storing, and processing data so it can be used for analysis, machine learning, and decision making.
- It focuses on:
 - data flow
 - system architecture
 - scalability
 - reliability
- It is infrastructure work, not analysis work.

Position in Data Science Ecosystem

- Typical ecosystem roles:
 - Data Engineer → prepares data
 - Data Scientist → analyzes and models
 - Data Analyst → interprets results
 - ML Engineer → deploys models
- Relationship:
Raw Data → Data Engineering → Prepared Data → Data Science
- Without this stage, analysis quality drops.

Goals of Data Engineering

- Accessibility
 - Data available when needed
- Consistency
 - Uniform format across sources
- Scalability
 - Handles growing volume
- Reliability
 - Pipelines run without failure
- Efficiency
 - Optimized resource use

Data Lifecycle Perspective

- Data engineering covers major stages of lifecycle:
 - Generation
 - Collection
 - Storage
 - Processing
 - Distribution
 - Archival or deletion

Data Storage Paradigms

- Relational storage
 - Usually structured tables which are schema enforced & SQL based
- NoSQL storage
 - Flexible schema stored as a document or key-value & scalable horizontally
- Data lakes
 - raw storage with large volume support
- Data warehouses
 - optimized for analytics
- Choice depends on workload.

Skills and Knowledge Areas

- Technical:
 - Programming
 - Database systems
 - Distributed computing
 - Networking basics
- Conceptual:
 - System design
 - Abstraction
 - Problem decomposition
-

Challenges in Practice

- Heterogeneous data sources
- Evolving schemas
- System bottlenecks
- Cost control
- Data privacy laws

Contents

6.1 Overview of data engineering in applied data science

6.2 Designing and implementing ETL pipelines

6.3 Automating workflows with schedulers (CRON, schedule)

6.4 Logging, monitoring, and error handling in pipelines

6.5 Data storage and retrieval strategies for pipelines

6.6 Automated report generation (Excel, HTML, PDF)

6.7 Case study: End-to-end automated analytics pipeline

ETL

- ETL (Extract, Transform, Load) is a structured process used to move data from source systems into storage systems prepared for analysis.
- It involves:
 - Extract — collecting raw data
 - Transform — cleaning and shaping
 - Load — storing in target system
- ETL ensures integration, consistency, and usability of data.

Purpose of ETL

- Combine multiple data sources
- Improve data quality
- Standardize formats
- Support analytics and reporting
- Enable historical storage
- ETL is a core mechanism in data warehouses and pipelines.

Extraction

- Extraction retrieves data from sources:
 - relational databases, APIs, logs, flat files, sensors etc.
- Challenges:
 - different formats
 - network latency
 - incomplete data
 - permission control

Methods for Extraction

- Full extraction
 - entire dataset copied
 - Simple but high resource cost
- Incremental extraction
 - only changed records
 - Efficient but requires tracking
- Change Data Capture (CDC)
 - detects real-time changes
 - used in modern pipelines

Transformation

- Transformation prepares data for use.
- Typical operations:
 - Cleaning
 - Normalization
 - Aggregation
 - Joining datasets
 - Filtering
 - Feature creation

Loading

- Loading transfers processed data to target system.
- Targets:
 - Databases
 - Warehouses
 - Data lakes
- Objectives:
 - Fast insertion
 - Integrity preservation
 - Indexing support

Loading Strategies

- Full load
 - replaces data
 - simple approach
- Incremental load
 - adds updates only
 - efficient
- Upsert strategy
 - update or insert
 - maintains current state

Contents

- 6.1 Overview of data engineering in applied data science
- 6.2 Designing and implementing ETL pipelines
- 6.3 Automating workflows with schedulers (CRON, schedule)**
- 6.4 Logging, monitoring, and error handling in pipelines
- 6.5 Data storage and retrieval strategies for pipelines
- 6.6 Automated report generation (Excel, HTML, PDF)
- 6.7 Case study: End-to-end automated analytics pipeline

Workflow Automation

- Workflow automation means running tasks automatically at planned times or events.
- In data pipelines this includes:
 - running ETL jobs
 - updating datasets
 - generating reports
 - backups
- Automation reduces manual work and errors.

Importance of Scheduling

- Ensures regular execution
- Maintains data freshness
- Supports reproducibility
- Reduces human intervention
- Improves reliability
- Without scheduling, pipelines depend on manual triggers.

Types of Scheduling

- Time-based
 - run at fixed times
- Event-based
 - run when condition met
- Dependency-based
 - run after another task finishes
- Modern systems often combine all three.

Introduction to CRON

- CRON is a time-based scheduler used in Unix/Linux systems.
- It executes commands automatically.
- Used for:
 - scripts
 - backups
 - ETL jobs
- Runs in background as a system service.

Introduction to CRON

- Cron is a software utility used to schedule jobs (commands or shell scripts) to run periodically at fixed times, dates, or intervals.
- **Cron Job:** A specific task scheduled to run automatically, such as daily backups or system maintenance.
- **Crontab:** The configuration file (cron table) where users define their schedules.

CRON Time Format

- A standard cron expression consists of five fields:
minute hour day_of_month month day_of_week
- Example pattern:
 - 0 0 * * * runs a command every night at midnight.
 - 30 2 * * * schedules a task to run at 2:30 AM every day.

Example

● ● ● -zsh

⌘1

```
pukarkarki@macbookpro ~ % EDITOR=nano crontab -e
```

Example

● ● ● nano

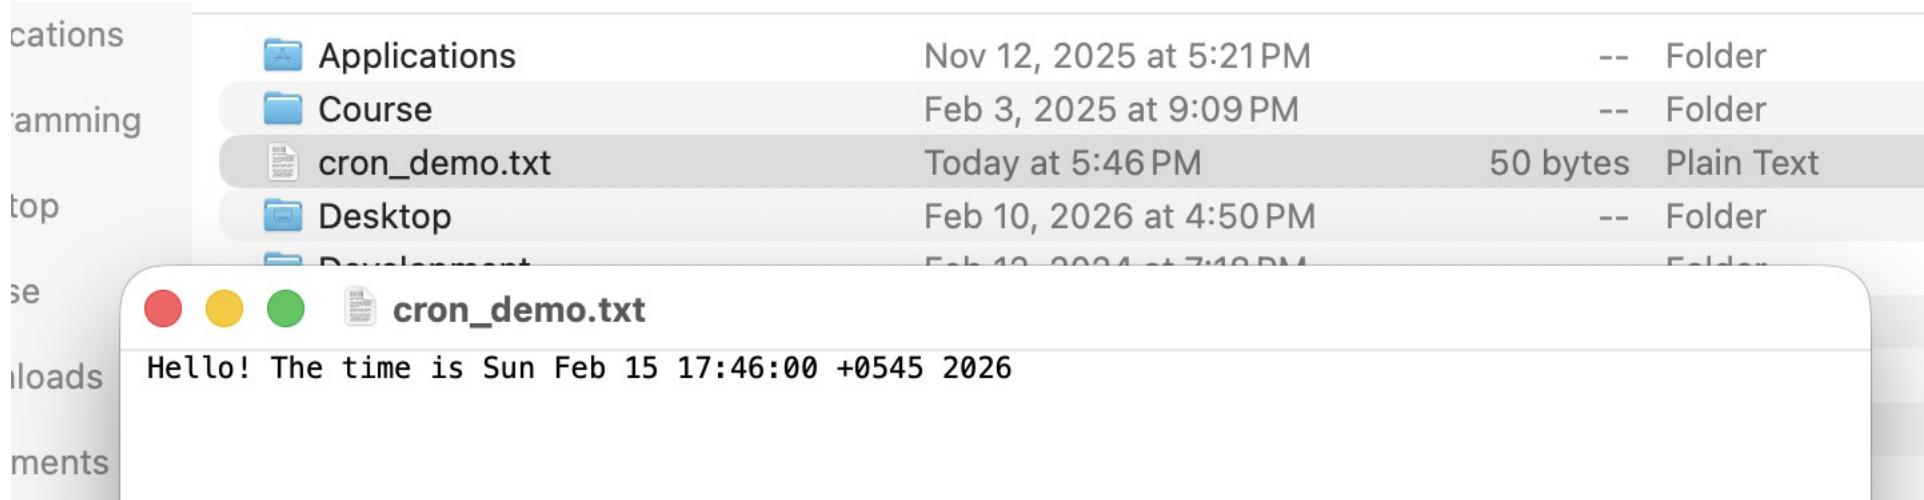
UW PICO 5.09

File: /tmp/crontab.m9cxLkMgPp

Modified

```
* * * * * echo "Hello! The time is $(date)" > /Users/pukarkarki/cron_demo.txt
```

Example



Example

● ● ● -zsh

29

```
pukarkarki@macbookpro ~ % crontab -l
* * * * * echo "Hello! The time is $(date)" > /Users/pukarkarki/cron_demo.txt
```

Real Life Examples of CRON

- **E-Commerce: Nightly Inventory Sync**
 - Goal: Sync stock levels between a physical warehouse and an online store.
 - Schedule: 0 0 * * * (Every night at midnight).
- **FinTech: Monthly Invoicing**
 - Goal: Generate and email PDF invoices to all active subscribers.
 - Schedule: 0 8 1 * * (8:00 AM on the 1st of every month).
- **Marketing: Daily Performance Reports**
 - Goal: Scrape website analytics and email a summary to the marketing team.
 - Schedule: 30 7 * * 1-5 (7:30 AM every weekday).

Advantages/Limitations of CRON

- Advantages
 - Lightweight
 - Reliable
 - Widely available
 - Simple configuration
- Limitations:
 - Limited dependency control
 - Basic monitoring
 - Manual logging setup needed

Python schedule Library

- **schedule** is a Python tool for simple job automation.
- Runs inside Python programs.
- Features:
 - readable syntax
 - easy integration
 - flexible intervals
- Best for small to medium tasks.

Python schedule Library

```
(AI) pukarkarki@macbookpro ~ % pip3 install schedule  
Collecting schedule  
  Downloading schedule-1.2.2-py3-none-any.whl.metadata (3.8 kB)  
  Downloading schedule-1.2.2-py3-none-any.whl (12 kB)  
Installing collected packages: schedule  
Successfully installed schedule-1.2.2
```

⚡ automate.py > ...

```
1 import schedule
2 import time
3 from datetime import datetime
4
5 def write_file():
6     with open("schedule_demo.txt", "w") as f:
7         f.write(f"Hello! The time is {datetime.now()}")
8
9 schedule.every(1).minutes.do(write_file)
10
11 while True:
12     schedule.run_pending()
13     time.sleep(1)
14
```

Python schedule Library

```
schedule.every(5).minutes.do(job)
schedule.every().hour.do(job)
schedule.every().day.at("14:30").do(job)
schedule.every().monday.do(job)
```

System Scheduler vs Application Scheduler

- System-level
 - CRON
 - independent of program
 - stable for servers
- Application-level
 - schedule library
 - embedded in code
 - flexible logic

Contents

- 6.1 Overview of data engineering in applied data science
- 6.2 Designing and implementing ETL pipelines
- 6.3 Automating workflows with schedulers (CRON, schedule)
- 6.4 Logging, monitoring, and error handling in pipelines**
- 6.5 Data storage and retrieval strategies for pipelines
- 6.6 Automated report generation (Excel, HTML, PDF)
- 6.7 Case study: End-to-end automated analytics pipeline

Logging, Monitoring, and Error Handling in Pipelines

- Data pipelines run automatically and often without direct supervision.
- To maintain reliability, systems must:
 - record activity
 - observe performance
 - handle failures
- This is achieved through:
 - Logging
 - Monitoring
 - Error handling
- These ensure stable pipeline operation.

Importance in Data Engineering

- Logging, Monitoring & Error handling are critical because
 - pipelines process large data
 - failures can go unnoticed
 - incorrect data affects analysis
 - downtime impacts decisions
- They provide visibility and control over pipeline behavior.

Logging

- Logging is the recording of system events during execution.
- Logs contain:
 - timestamps
 - task status
 - messages
 - error details

Types of Logs

- Execution logs
 - track job progress
- Error logs
 - record failures
- Audit logs
 - track user/system actions
- Performance logs
 - capture timing and usage

Logging Levels

- Standard severity levels:
 - DEBUG — detailed internal info
 - INFO — normal operation
 - WARNING — potential issues
 - ERROR — failures occurred
 - CRITICAL — system stopping issues

Logging Best Practices

- Use structured messages
- Include timestamps
- Log meaningful context
- Avoid excessive verbosity
- Store logs centrally

Good logs simplify debugging.

Monitoring

- Monitoring is continuous observation of system health and performance.
- It answers:
 - Is pipeline running?
 - How fast?
 - Any failures?
- Provides real-time awareness.

Monitoring Metrics

- Common metrics:
 - Execution time
 - Success rate
 - Throughput
 - Resource usage
 - Queue length
- Metrics indicate pipeline efficiency.

Monitoring Methods

- Dashboards
- Automated alerts
- Periodic checks
- Visual analytics

Error Handling

- Error handling is the structured response to unexpected issues during execution.
- Goals:
 - prevent crashes
 - protect data integrity
 - recover gracefully
- This maintains pipeline continuity.

Types of Errors

- Syntax errors
 - Coding mistakes
- Runtime errors
 - Occur during execution
- Logical errors
 - Incorrect results
- System errors
 - Hardware/network failure

Techniques for Error Handling

- Try–except structures
- Retries
- Fallback steps
- Checkpoint recovery
- Transaction rollback

Fault Tolerance

- Fault tolerance means continuing operation despite failures.
- Achieved through:
 - Redundancy
 - Replication
 - Distributed execution
 - Checkpoint storage
- Important for large-scale systems.

Contents

- 6.1 Overview of data engineering in applied data science
- 6.2 Designing and implementing ETL pipelines
- 6.3 Automating workflows with schedulers (CRON, schedule)
- 6.4 Logging, monitoring, and error handling in pipelines
- 6.5 Data storage and retrieval strategies for pipelines**
- 6.6 Automated report generation (Excel, HTML, PDF)
- 6.7 Case study: End-to-end automated analytics pipeline

Data Storage and Retrieval Strategies for Pipelines

- Data pipelines generate and process large volumes of data.
- Efficient storage and retrieval strategies ensure:
 - fast access
 - reliability
 - scalability
 - cost control
- This stage determines overall pipeline performance.

Goals of Storage Strategy

- Preserve data integrity
- Support efficient querying
- Enable scalability
- Ensure availability
- Optimize cost

Different Storage Architecture

- Data moves through layers:
Raw Storage → Processed Storage → Analytical Storage
- Each layer serves a different purpose and format.

Types of Storage

- Relational Databases
 - structured tables
 - strong consistency
- NoSQL Databases
 - flexible schema
 - high scalability
- Data Warehouses
 - analytics optimized

Relational Storage

- Schema-based structure
- SQL querying
- Transactional integrity
- Normalization support
- Best suited for structured data and reporting.
- Limitations:
 - less flexible scaling
 - rigid schema changes

NoSQL Storage

- Types:
 - document stores
 - key-value stores
 - column-based
 - graph databases
- Advantages:
 - flexible schema
 - horizontal scaling
- Trade-off:
 - weaker consistency guarantees

Date Warehouse

- Designed for analytics.
- Features:
 - integrated data sources
 - historical storage
 - optimized aggregation queries
 - columnar storage models
- Supports decision making and reporting.

Contents

- 6.1 Overview of data engineering in applied data science
- 6.2 Designing and implementing ETL pipelines
- 6.3 Automating workflows with schedulers (CRON, schedule)
- 6.4 Logging, monitoring, and error handling in pipelines
- 6.5 Data storage and retrieval strategies for pipelines
- 6.6 Automated report generation (Excel, HTML, PDF)**
- 6.7 Case study: End-to-end automated analytics pipeline

Automated Report Generation

- Automated report generation is the process of producing formatted outputs from processed data without manual intervention.
- Reports communicate results to users in readable form.
- Common formats:
 - Excel
 - HTML
 - PDF
- This is usually the final stage of a pipeline.

Automated Report Generation

- Why?
 - Deliver insights regularly
 - Reduce manual work
 - Ensure consistency
 - Improve decision support
 - Support large-scale distribution

Reporting Formats

- Excel
 - Interactive tables
 - User manipulation possible
- HTML
 - Web display
 - Accessible remotely
- PDF
 - Fixed layout
 - Professional sharing
- Choice depends on audience needs.

Automation Workflow

- **Steps:**
 - 1) Data analysis completed
 - 2) Format template selected
 - 3) Content inserted
 - 4) Report exported
 - 5) Delivered to users
- Reports often generated daily, weekly or monthly and schedulers trigger generation automatically after pipeline completion.

Contents

- 6.1 Overview of data engineering in applied data science
- 6.2 Designing and implementing ETL pipelines
- 6.3 Automating workflows with schedulers (CRON, schedule)
- 6.4 Logging, monitoring, and error handling in pipelines
- 6.5 Data storage and retrieval strategies for pipelines
- 6.6 Automated report generation (Excel, HTML, PDF)
- 6.7 Case study: End-to-end automated analytics pipeline**

Q) Design an end-to-end automated analytics pipeline for a real-world organization (such as retail, hospital, banking, or education). Explain:

- data sources
- ETL process
- storage choice
- scheduling method
- logging and monitoring
- report generation